

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: vaccinetrain=pd.read_csv("training_set_features.csv")
```

```
In [3]: vaccinetest=pd.read_csv("test_set_features.csv")
```

```
In [4]: vaccine_labels=pd.read_csv("training_set_labels.csv")
```

```
In [5]: print(vaccinetrain.shape)
print(vaccinetest.shape)
print(vaccine_labels.shape)
```

(26707, 36)

(26708, 36)

(26707, 3)

```
In [6]: vaccinetrain
```

```
Out[6]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoida
0	0	1.0	0.0	0.0	
1	1	3.0	2.0	0.0	
2	2	1.0	1.0	0.0	
3	3	1.0	1.0	0.0	
4	4	2.0	1.0	0.0	
...	
26702	26702	2.0	0.0	0.0	
26703	26703	1.0	2.0	0.0	
26704	26704	2.0	2.0	0.0	
26705	26705	1.0	1.0	0.0	
26706	26706	0.0	0.0	0.0	

26707 rows × 36 columns

In [7]: vaccinetest

Out[7]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoida
0	26707	2.0	2.0	0.0	
1	26708	1.0	1.0	0.0	
2	26709	2.0	2.0	0.0	
3	26710	1.0	1.0	0.0	
4	26711	3.0	1.0	1.0	
...
26703	53410	1.0	1.0	0.0	
26704	53411	3.0	1.0	0.0	
26705	53412	0.0	1.0	0.0	
26706	53413	3.0	1.0	0.0	
26707	53414	2.0	1.0	0.0	

26708 rows × 36 columns



In [8]: vaccinetrain['id']='train'
vaccinetest['id']='test'

In [9]: combinedf=pd.concat([vaccinetrain,vaccinetest],axis=0)

In [10]: combinedf.shape

Out[10]: (53415, 37)

```
In [11]: combinedf.isnull().sum().sort_values(ascending=False)
```

```
Out[11]: employment_occupation      26896
employment_industry      26605
health_insurance      24502
income_poverty      8920
doctor_recc_h1n1      4320
doctor_recc_seasonal      4320
rent_or_own      4078
employment_status      2934
marital_status      2850
education      2814
chronic_med_condition      1903
child_under_6_months      1633
health_worker      1593
opinion_seas_sick_from_vacc      1058
opinion_seas_risk      1013
opinion_seas_vacc_effective      914
opinion_h1n1_vacc_effective      789
opinion_h1n1_sick_from_vacc      770
opinion_h1n1_risk      768
household_children      474
household_adults      474
behavioral_avoidance      421
behavioral_touch_face      256
h1n1_knowledge      238
h1n1_concern      177
behavioral_outside_home      164
behavioral_large_gatherings      159
behavioral_antiviral_meds      150
behavioral_wash_hands      82
behavioral_face_mask      38
census_msa      0
respondent_id      0
hhs_geo_region      0
sex      0
race      0
age_group      0
id      0
dtype: int64
```

```
In [12]: for col in combinedf.drop(['respondent_id', 'id'], axis=1).columns:
         freq=combinedf[col].value_counts(dropna=False)
         print(freq)
```

```
2.0    21318
1.0    16311
3.0     9114
0.0     6495
NaN       177
Name: h1n1_concern, dtype: int64
1.0    29227
2.0    19002
0.0     4948
NaN       238
Name: h1n1_knowledge, dtype: int64
0.0    50642
1.0     2623
NaN       150
Name: behavioral_antiviral_meds, dtype: int64
1.0    38564
0.0    14430
NaN       421
Name: behavioral_avoidance, dtype: int64
0.0    40607
```

```
In [13]: #non-parametric algorithms means algorithms which has no prior assumptions
         #ex : KNN (K-nearest neighbours)
         #missing values imputation (KNNImputer)
         #Label encoding must be done before imputation
         #Label encoder will keep the missing values as it is
```

```
In [14]: from sklearn.preprocessing import LabelEncoder
```

```
In [15]: stringcols=combinedf.drop(['respondent_id', 'id'], axis=1)
```

```
In [16]: stringcols=stringcols.astype(str).apply(lambda series:pd.Series(
         LabelEncoder().fit_transform(series[series.notnull()]),
         index=series[series.notnull()].index))
```

```
In [17]: from sklearn.impute import KNNImputer
```

```
In [18]: imputer=KNNImputer()
```

```
In [19]: stringcolsimp=imputer.fit_transform(stringcols)
```

```
In [20]: stringcols.isnull().sum().sort_values(ascending=False)
```

```
Out[20]: h1n1_concern          0
marital_status          0
opinion_seas_sick_from_vacc  0
age_group              0
education              0
race                  0
sex                  0
income_poverty        0
rent_or_own          0
opinion_seas_vacc_effective  0
employment_status     0
hhs_geo_region        0
census_msa           0
household_adults      0
household_children    0
employment_industry   0
opinion_seas_risk      0
opinion_h1n1_sick_from_vacc  0
h1n1_knowledge        0
behavioral_touch_face  0
behavioral_antiviral_meds  0
behavioral_avoidance   0
behavioral_face_mask   0
behavioral_wash_hands  0
behavioral_large_gatherings  0
behavioral_outside_home  0
doctor_recc_h1n1      0
opinion_h1n1_risk      0
doctor_recc_seasonal  0
chronic_med_condition  0
child_under_6_months  0
health_worker         0
health_insurance      0
opinion_h1n1_vacc_effective  0
employment_occupation  0
dtype: int64
```

```
In [21]: stringcolsimp=pd.DataFrame(stringcolsimp,columns=stringcols.columns)
```

In [22]: stringcolsimp

Out[22]:

	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance	behavioral_
0	1.0	0.0	0.0	0.0	
1	3.0	2.0	0.0	1.0	
2	1.0	1.0	0.0	1.0	
3	1.0	1.0	0.0	1.0	
4	2.0	1.0	0.0	1.0	
...
53410	1.0	1.0	0.0	1.0	
53411	3.0	1.0	0.0	1.0	
53412	0.0	1.0	0.0	0.0	
53413	3.0	1.0	0.0	1.0	
53414	2.0	1.0	0.0	0.0	

53415 rows × 35 columns

In [23]:

```
for col in stringcolsimp.columns:
    freq=stringcolsimp[col].value_counts(dropna=False)
    print(freq)
```

```
2.0    21318
1.0    16311
3.0     9114
0.0     6495
4.0       177
Name: h1n1_concern, dtype: int64
1.0    29227
2.0    19002
0.0     4948
3.0       238
Name: h1n1_knowledge, dtype: int64
0.0    50642
1.0     2623
2.0       150
Name: behavioral_antiviral_meds, dtype: int64
1.0    38564
0.0    14430
2.0       421
Name: behavioral_avoidance, dtype: int64
0.0    10607
```

In [24]: vaccinetraindf=stringcolsimp.loc[0:26706]

In [25]: vaccinetestdf=stringcolsimp.loc[26707:53415]

```
In [26]: print(vaccinetraindf.shape)
print(vaccinetestdf.shape)
```

```
(26707, 35)
(26708, 35)
```

LogisticRegression

```
In [27]: from sklearn.linear_model import LogisticRegression
```

```
In [28]: logreg=LogisticRegression()
```

```
In [29]: X=vaccinetraindf
y=vaccine_labels.h1n1_vaccine
```

```
In [30]: logregmodel=logreg.fit(X,y)
```

```
C:\Users\admin\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:76
3: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
In [31]: logregmodel.score(X,y)
```

```
Out[31]: 0.8280975025274273
```

```
In [32]: logregpredict=logregmodel.predict(X)
```

```
In [33]: from sklearn.metrics import classification_report
```

```
In [34]: print(classification_report(y,logregpredict))
```

	precision	recall	f1-score	support
0	0.85	0.95	0.90	21033
1	0.67	0.37	0.48	5674
accuracy			0.83	26707
macro avg	0.76	0.66	0.69	26707
weighted avg	0.81	0.83	0.81	26707

```
In [35]: y.value_counts()
```

```
Out[35]: 0    21033
         1     5674
         Name: h1n1_vaccine, dtype: int64
```

CategoricalNaiveBayes

```
In [36]: from sklearn.naive_bayes import CategoricalNB
```

```
In [37]: catnb=CategoricalNB()
```

```
In [38]: catnb=CategoricalNB()
```

```
In [39]: catNBmodel.score(X,y)
```

```
Out[39]: 0.8092260456060208
```


```
In [40]: catNBpredict=catNBmodel.predict(X)
```

```
In [41]: print(classification_report(y,catNBpredict))
```

	precision	recall	f1-score	support
0	0.89	0.86	0.88	21033
1	0.55	0.61	0.58	5674
accuracy			0.81	26707
macro avg	0.72	0.74	0.73	26707
weighted avg	0.82	0.81	0.81	26707



```
In [42]: pd.value_counts(y).plot(kind="pie", autopct="%f")
```

```
Out[42]: <AxesSubplot:ylabel='h1n1_vaccine'>
```



```
In [43]: pd.crosstab(X.h1n1_concern,y).plot(kind='bar',stacked=False)
```

```
Out[43]: <AxesSubplot:xlabel='h1n1_concern'>
```



```
In [44]: combinedf.columns
```

```
Out[44]: Index(['respondent_id', 'h1n1_concern', 'h1n1_knowledge',
                'behavioral_antiviral_meds', 'behavioral_avoidance',
                'behavioral_face_mask', 'behavioral_wash_hands',
                'behavioral_large_gatherings', 'behavioral_outside_home',
                'behavioral_touch_face', 'doctor_recc_h1n1', 'doctor_recc_seasonal',
                'chronic_med_condition', 'child_under_6_months', 'health_worker',
                'health_insurance', 'opinion_h1n1_vacc_effective', 'opinion_h1n1_risk',
                'opinion_h1n1_sick_from_vacc', 'opinion_seas_vacc_effective',
                'opinion_seas_risk', 'opinion_seas_sick_from_vacc', 'age_group',
                'education', 'race', 'sex', 'income_poverty', 'marital_status',
                'rent_or_own', 'employment_status', 'hhs_geo_region', 'census_msa',
                'household_adults', 'household_children', 'employment_industry',
                'employment_occupation', 'id'],
                dtype='object')
```

```
In [45]: combinedf.head()
```

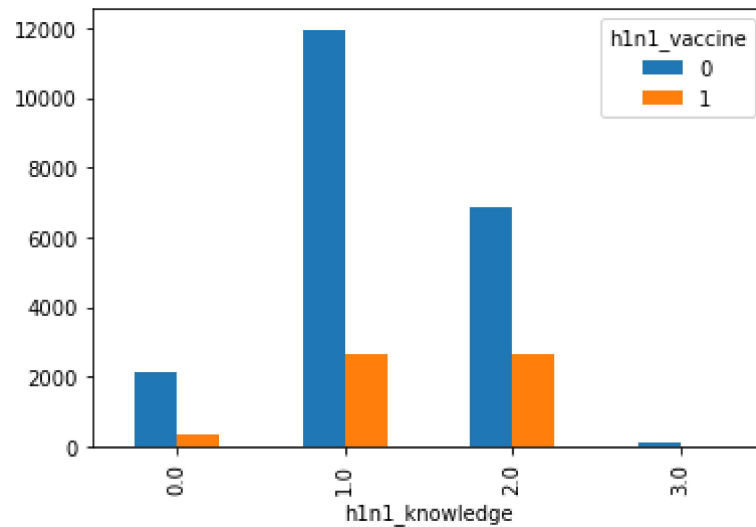
```
Out[45]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance
0	0	1.0	0.0	0.0	0.0
1	1	3.0	2.0	0.0	1.0
2	2	1.0	1.0	0.0	1.0
3	3	1.0	1.0	0.0	1.0
4	4	2.0	1.0	0.0	1.0

5 rows × 37 columns


```
In [46]: pd.crosstab(X.h1n1_knowledge,y).plot(kind='bar',stacked=False)
```

```
Out[46]: <AxesSubplot:xlabel='h1n1_knowledge'>
```



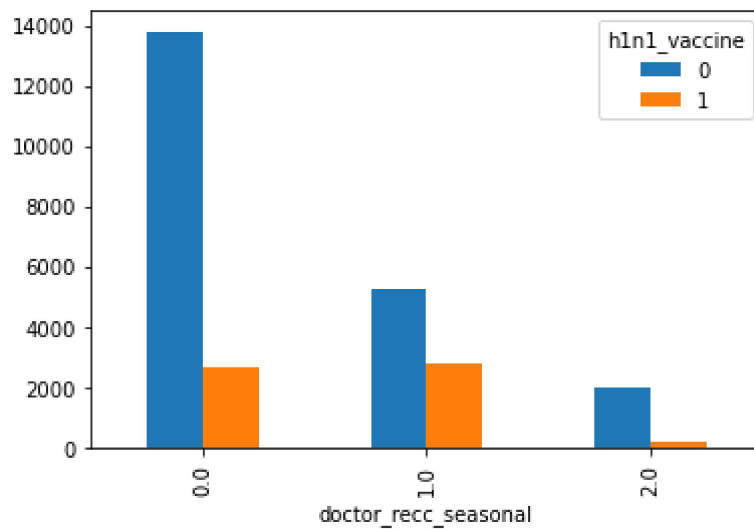
```
In [47]: pd.crosstab(X.doctor_recc_h1n1,y).plot(kind='bar',stacked=False)
```

```
Out[47]: <AxesSubplot:xlabel='doctor_recc_h1n1'>
```



```
In [48]: pd.crosstab(X.doctor_recc_seasonal,y).plot(kind='bar',stacked=False)
```

```
Out[48]: <AxesSubplot:xlabel='doctor_recc_seasonal'>
```



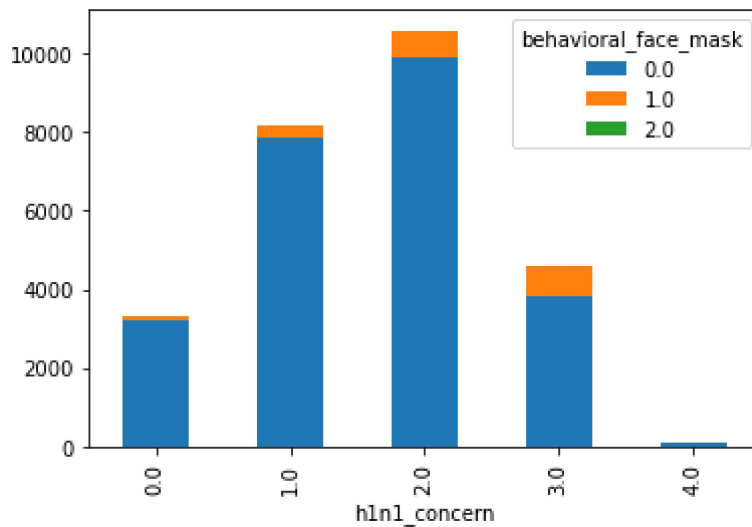
```
In [49]: pd.crosstab(X.h1n1_concern,y)
```

```
Out[49]:
```

	h1n1_vaccine 0	h1n1_vaccine 1
h1n1_concern		
0.0	2849	447
1.0	6756	1397
2.0	8102	2473
3.0	3250	1341
4.0	76	16

```
In [50]: pd.crosstab(X.h1n1_concern,X.behavioral_face_mask).plot(kind='bar',stacked=True)
```

```
Out[50]: <AxesSubplot:xlabel='h1n1_concern'>
```



MultiLevelPerceptron

```
In [51]: from sklearn.neural_network import MLPClassifier #Multi-Level Perceptron()
```

```
In [54]: nn=MLPClassifier(hidden_layer_sizes=(50,100,50),activation='logistic',max_iter=1000)
```

```
In [55]: nnmodel=nn.fit(X,y)
```

```
In [56]: nnmodel.score(X,y) #h1n1 vaccine
```

```
Out[56]: 0.8757629086007414
```